

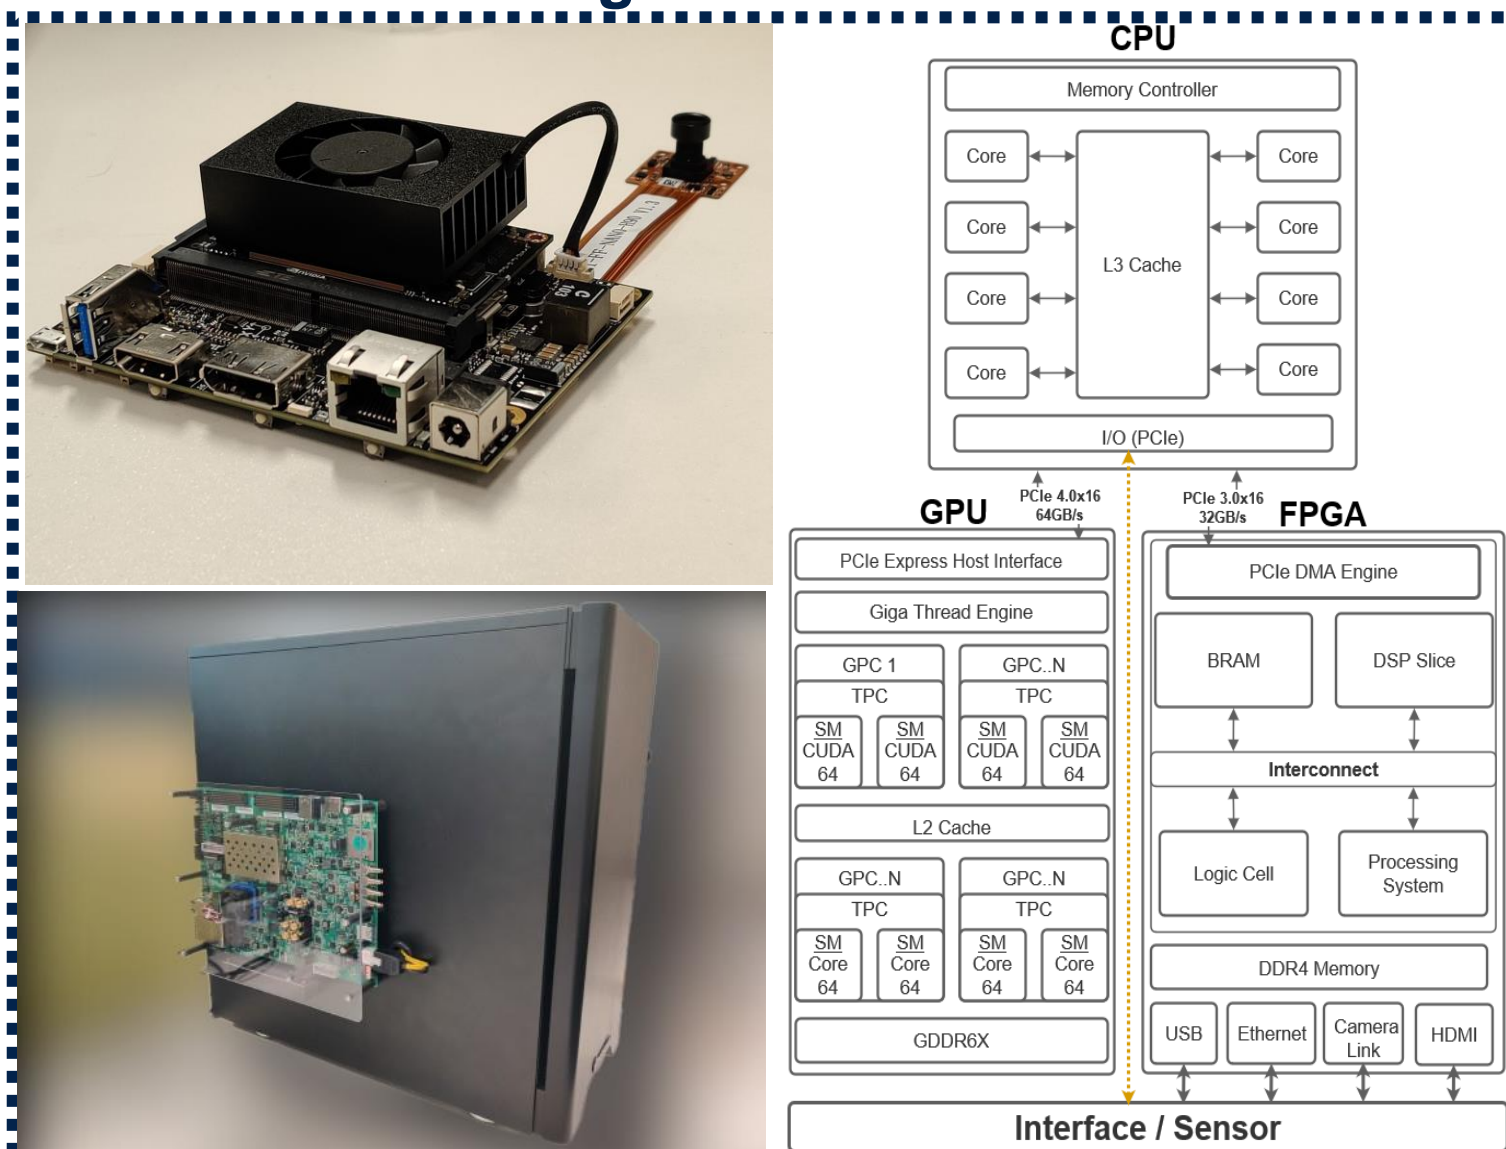
# Domain-specific Optimisations of Image Processing Algorithms on Heterogenous Architectures.

Teymoor Ali<sup>1,2,3</sup>, Deepayan Bhowmik<sup>2</sup>, Robert Nicol<sup>3</sup>, University of Stirling<sup>1</sup>/Newcastle University<sup>2</sup>/STMicroelectronics<sup>3</sup>

Email: [teymoor.ali@newcastle.ac.uk](mailto:teymoor.ali@newcastle.ac.uk), [deepayan.bhowmik@newcastle.ac.uk](mailto:deepayan.bhowmik@newcastle.ac.uk), [robert.nicol@st.com](mailto:robert.nicol@st.com)

**Heterogeneous computing architectures containing CPUs, GPUs, FPGAs and other specialised accelerators on a single interconnected fabric have enabled widespread adoption within many real-time vision applications.**

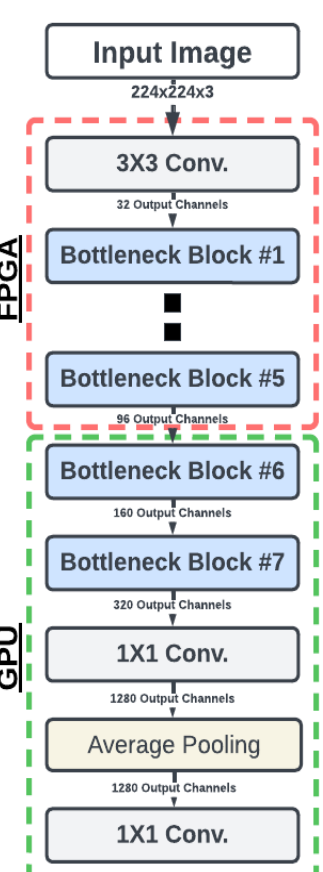
## Heterogenous Platforms



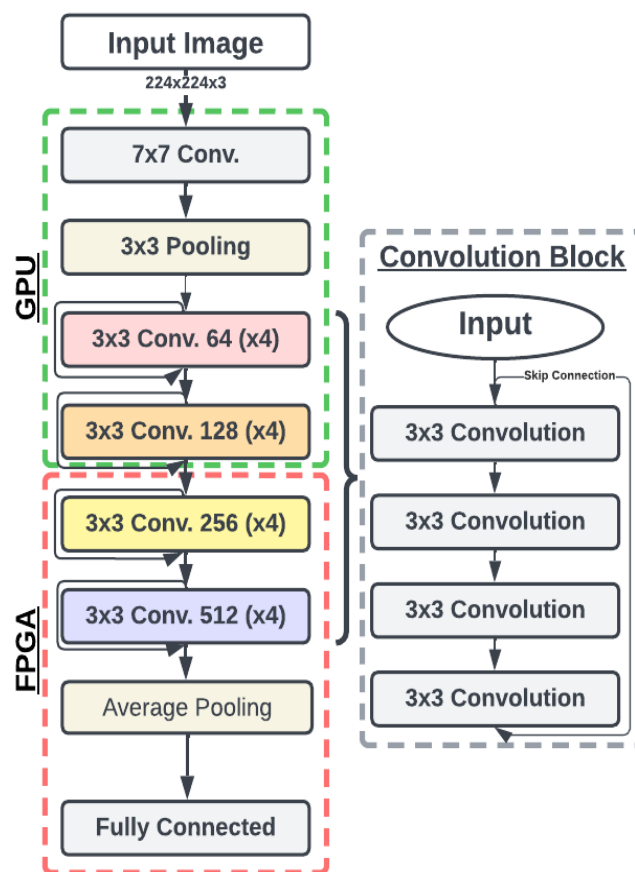
## Partitioning Strategies

- Designing algorithms on heterogenous architectures requires careful partitioning of operations.

### MobileNetV2




### ResNet-18



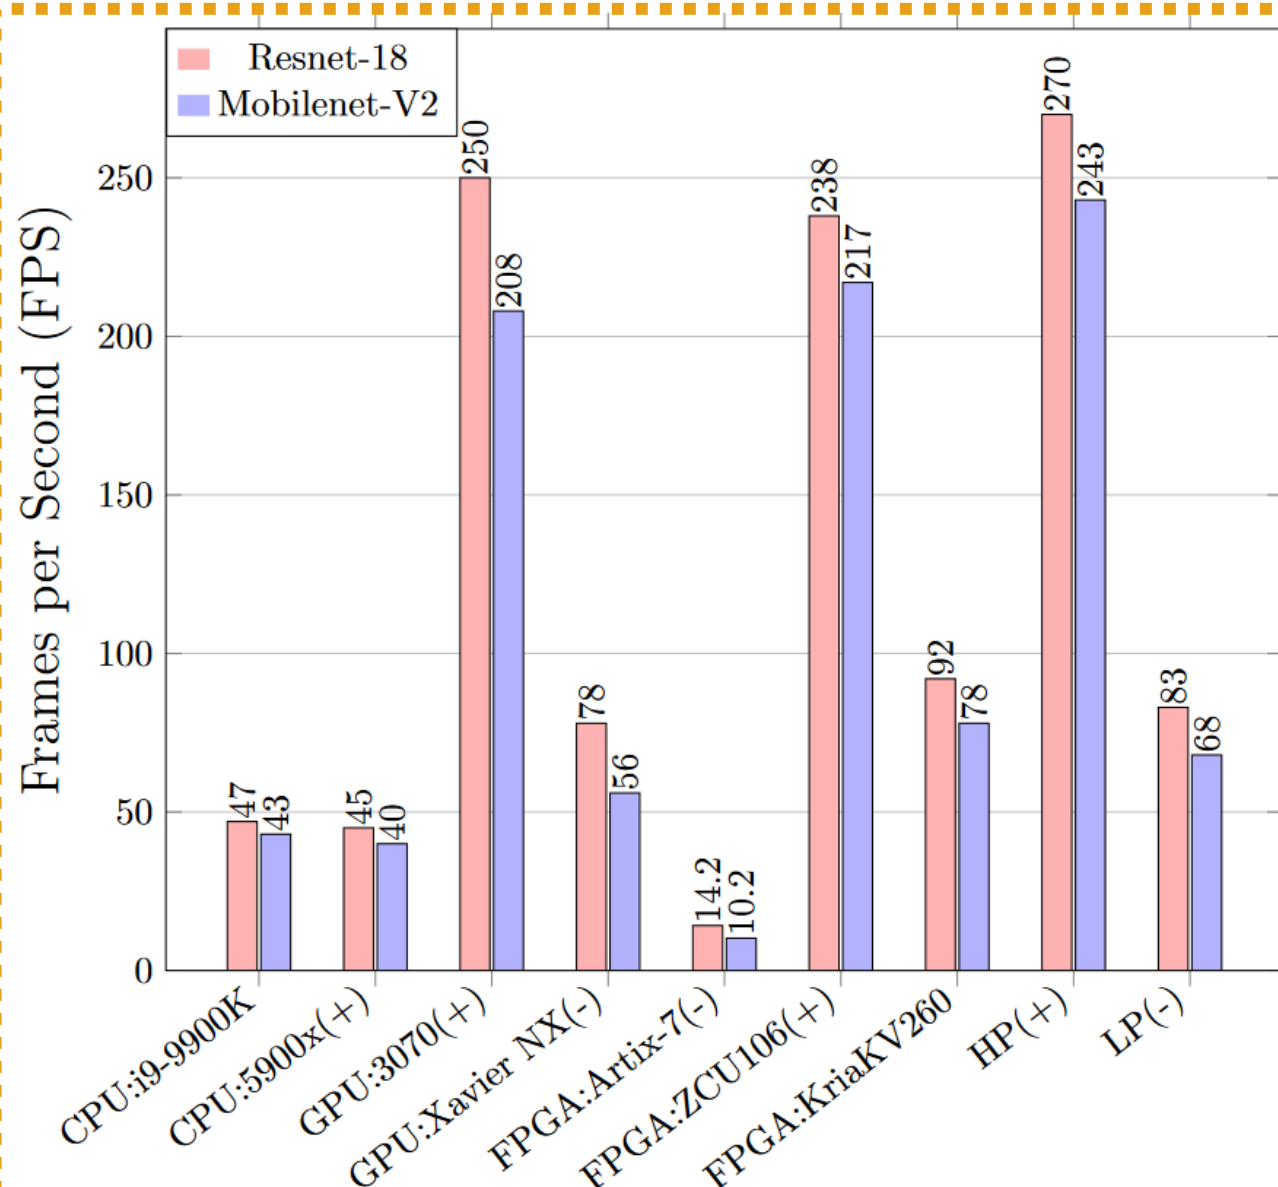
## Domain-Specific Optimisations

Separable Filter:  $\frac{1}{4} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} = \frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$



Kernel Size:  $\frac{1}{110} \begin{bmatrix} 1 & 1 & 2 & 2 & 2 & 1 & 1 \\ 1 & 2 & 2 & 4 & 2 & 2 & 1 \\ 2 & 2 & 4 & 8 & 4 & 2 & 2 \\ 2 & 4 & 8 & 16 & 8 & 4 & 2 \\ 2 & 2 & 4 & 8 & 4 & 2 & 2 \\ 1 & 2 & 2 & 4 & 2 & 2 & 1 \\ 1 & 1 & 2 & 2 & 2 & 1 & 1 \end{bmatrix} \rightarrow \frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$

## Benchmarking CNNs on Hardware



## Research Outcomes

- Domain-specific optimisations techniques applied on imaging algorithms, significantly reduce execution time and energy consumption of hardware.
- Partitioning algorithms with respect to memory latency onto the most efficient accelerator in a heterogenous domain increases overall performance.

## Benchmarking Results

- High and low-power heterogenous platforms outperforms their fastest discrete homogenous counterpart architecture (3070)/(XavierNX) in terms of execution time and energy.
- Kria FPGA outperforms in FPS and energy than embedded GPU counterpart (XavierNX) at a lower clock speed.